Does It Scale? Static Output Approximation of PHP Web Applications Using Symbolic Execution[*]

Stefan Mühlbauer
Technische Universität Braunschweig

Christian Kästner
Carnegie Mellon University

Tien N. Nguyen
University of Texas at Dallas

ABSTRACT

Dynamic web applications have become widely popular and are to a large proportion based on the scripting language PHP. Output approximation of web applications enables a range of additional tool support as well as possibilities for vulnerability detection. Unfortunately, recent approximation approaches have only been evaluated for smaller systems.

This paper presents an experience report about the scalability of output approximation using symbolic execution of state-of-the-art PHP web applications. For a symbolic execution engine extended with support for object-oriented programming and arrays, we identified language features and corresponding programming patterns that impede symbolic execution and limit the scalability of this approach. Our findings include: (1) Dynamic features such as functions and includes are prone to fail for certain programming patterns. (2) Expressions containing elements from I/O, databases or files can heavily impede symbolic execution. Our findings provide useful guidelines to design new tools and also to improve the development process of statically analyzable web applications.

CCS CONCEPTS

•**Computer systems organization Embedded systems;** *Redundancy;* Robotics; •**Networks** Network reliability;

KEYWORDS

ACM proceedings, LaTeX, text tagging

1 INTRODUCTION

With the emerging world wide web, dynamic web applications have become widely popular. Various different implementation techniques have emerged, ranging from technologies for programming languages (JSP, ASP .NET) and web application frameworks for script languages (Ruby On Rails, Django) to languages tailored specifically to the domain of web applications. PHP citephpNET is

---

[*]Produces the permission block, and copyright information

a programming language focused on server-side application development. As of 2012, it was used by 78.8 percent of the ten million most popular websites (according to Alexa popularity ranking) [3], ranking 7th on the TIOBE programming community index [4], and was the ranked as the 6th most popular language on GitHub [2]. One common property of all technologies for dynamic web applications is *staged computation*: A dynamic web application as a whole consists of both static code, such as scripts, and dynamically generated code, such as client page output. The latter code, although assembled at runtime, may contain client-side parts of the web application, such as JavaScript. So, to study dynamic web applications in its entirety, we need to consider both static as well as dynamic aspects of systems.

Symbolic execution is a static analysis technique between program testing and program proving. For a program, symbolic values are used instead of concrete inputs [5, 7]. The underlying concept is to map program input to program output: Symbolic values are used and propagated throughout the symbolic execution and keep dependencies for program output traceable. This analysis technique helps unfold the *staged* nature of dynamic web applications. Since every feasible path can be executed the symbolic output contains both invariant output as well as output variants that depend on program input.

Knowledge about different output output variants is leveraged by a number of analyses for the domain of web applications. Previous work presented tools for detecting and locating HTML validation errors [12], computing program slices across server-side and client-page code [10], or easing development and maintenance by extending IDE support for web applications with code navigation [9, 11].

Despite the various use cases for analyses based on an approximated output model, static output approximation for PHP web applications using symbolic execution has so far only been evaluated for small systems that are not maintained anymore. The tools presented by previous work [9–12] though are only practical, if for a given system the symbolic execution engine is scalable, i.e., will also approximate output accurately for larger and more recent systems with acceptable time and space consumption.

To investigate the question, whether we can have a practical and scalable symbolic executoin engine, we re-implemented the engine with the specifications of the previous symbolic execution semantics [9] for PHP and additional features, such as object-oriented programming. We evaluate our symbolic execution engine for large-scale and modern PHP systems.

During the introduction of new semantics, we addressed two tradeoffs between accuracy of our e xecution and performance: For method calls with ambiguous targets as well as concrete execution of loops the number of program states can become infeasible. This

may require additional engineering adjustments to tame the state space explosion at cost of accuracy.

Based on out observations we identify conceptual limitations of symbolic execution for PHP. Dynamic features such as indirect function calls and include expression require concrete information to be evaluated properly. Since these expressions may be assembled at runtime and can contain symbolic information, for many cases this restricts our engine from further execution. Although some effort to statically approximate include expression can be spent, it is often non-trivial since expressions can contain information coming from various inputs including databases, user inputs or configuration files. We empirically evaluate our observations for a corpus including large PHP systems that are still maintained.

Our key contribution in this paper include (1) a new tool infrastructure to statically approximate the client-page output for PHP web applications using symbolic execution and (2) a report of out observations for state-of-the-art web applications as well as an empirical evaluation of conceptual limitations of symbolic executions for PHP web applications.

2 STATE OF THE ART

Symbolic execution is an analysis techniques used by the tools whose practicality we aim to evaluate. This section recaps the idea of symbolic execution, the tools which make use of it, and both conceptual and language-related limitations.

2.1 Symbolic Execution

Symbolic execution is a static analysis technique that was proposed by James C. King in 1976 [7]. Symbolic execution allows to explore feasible paths in a program. Based on normal program execution, the execution semantics is extended to handle symbolic values. An execution becomes symbolic by assuming symbolic values as program inputs rather than obtaining concrete values [5, 7]. These symbolic values are then propagated and used throughout the execution.

Execution starts with a tautology path condition. The control-flow can be split whenever the decision between branches is ambiguous. This is the case, when conditional expressions, for example for if-statements, evaluate to a symbolic value rather than an actual result. The corresponding path conditions (guards for instance) are conjoined with the previous path condition [7]. Additionally, the conjoined path condition can be checked for satisfiability in order to exclude infeasible paths.

2.2 Static Output Approximation for PHP

Symbolic execution is not dedicated to a particular programming language. For dynamic web applications, in our case for web applications written in PHP, it can be used to approximate all possible client page output variants. Before we present in detail existing tools that aim to ease development and maintenance of PHP web applications, we illustrate how output approximation using symbolic execution is pursued by starting with a small example.

Figure 1 shows two code listings. The snippet of server-side web application code in Figure 1a contains both HTML code (input form in

```
1  <form>
2    <input type="text" name="name">
3    <input type="submit" value="Submit">
4  </form>
5
6  <?php
7  function make_titles($arg, $n) {
8    $is = shuffle(array(1,2,3));
9    while ($n < $is[0]) {
10     echo "<h1>" . $arg . "</h1><br />";
11     $n = $n + 1;
12   }
13 }
14
15 call_user_func("make_titles", "Headline", 4);
16
17 $greeting = "Hello " . $_POST['name'] .  "!";
18 if (!isset($_POST['name'])) {
19   echo "No name entered!";
20 } else {
21   echo $greeting;
22 }
23 ?>
```

(a) Snippet of HTML and PHP server-side web application code.

```
1  <form>
2    <input type="text" name="name">
3    <input type="submit" value="Submit">
4  </form>
5
6  // #repeat n < shuffle(is)
7  <h1>Headline</h1><br />
8  // #endrepeat
9
10 // #if !isset(name)
11 No name entered!
12 // #else
13 Entered $\alpha$!
14 // #endif
```

(b) Approximation client paege output of the server-side web application code in Figure 1a. Variability as well as repetition annotations are represented by preprocessor directives.

Figure 1: Illustration of static ouptut approximation using symbolic execution

lines 1 – 4) as well as PHP code: First, the function `make_titles` defined in lines 7 – 13 prints the string `$arg` passed as an argument between zero and an arbitrary number of times. Second, the function is indirectly called using the built-in mechnism `call_user_func` in line 15 with the string `Headline` and 4 as arguments. Finally, the input from the form in lines 1 – 4 is used: The variable `$greeting` consists of three concatenated strings, `Hello`, the value of input field `name`, and an exclamation mark. For the greeting there are two possible client page output variants: If the input field `name` is empty at runtime, line 19 is executed and `No name entered!` is printed, or, if a name was entered, line 21 is executed and the greeting constructed in line 17 is printed.

The following symbolic execution semantics for PHP features based on previous work [9] illustrate the necessary modifications and extensions made to handle the level of ambiguity that comes with handling symbolic values. In addition to normal output, the output

approximation contains preprocessor directives to express repeatable parts as well as output variants. All approximated client page output corresponding to Figure 1a consisting of two different variants is illustrated in Figure 1b. First, for loops in the previous execution semantics [9] only one iteration is executed instead of an arbitrary number of iterations. Since a loop condition can contain symbolic values, it might not be feasible to determine whether, and if so, when a loop execution terminates. Hence, any output constructed during the single loop iteration is highlighted (`repeat ...endrepeat`) to possibly be repeated an arbitrary number of times. The output in line 7 is marked to be repeated an arbitrary number of times, as the number of loop iterations cannot be determined statically. Second, a similar procedure applies to recursive functions. Once called, a recursive function is also only executed to recursion depth one per call, and any subsequent invocation of that function returns a symbolic value. This, again, is due to the difficulty to determine the exact recursion depth similar to the number of loop iterations. Third, any form of user interaction providing input data as well as any form of input referred to from the deployment context of the web application, i.e. configuration files or databases, is assumed to be symbolic. The context variable `$_POST` represents input submitted to the server and, hence, is symbolic as every program input is per definition substituted by symbolic values. Finally, the output in lines 11 and 13 respectively are marked as alternative variants (`if ...else ...endif`) as symbolic execution of lines 18 – 22 in Figure 1a will explore both different branches.

2.3 Existing Tool Support for PHP

All tools proposed so far leverage a representation of all possible HTML client page output. Since any single client page output variant can be analyzed this enables tool support addressing the web application as a whole in spite of its staged nature, where parts of the web application are generated dynamically. Based on the approximated output representation, subsequent analyses can be conducted:

- *PhpSync*: Using all variants of HTML client page output, every single one can be statically checked for markup validity, i.e., if the web page conforms to syntactical specifications for HTML and other client-side languages. After the tool detects validation errors, auto-fixes can be provided. Otherwise, it traces validation errors back to source code responsible for the defect, which then can be refactored manually [12].
- *WebSlice*: Program slices enable to extract and understand the impact of changes in an application. To consider all of a dynamic web application for program slices, client page output needs to be taken into account. WebSlice combines PHP data-flow information with an output approximation to enable program slices across different languages [10].
- *Varis*: Editor services such as "jump to declaration" are nontrivial for dynamic web applications due to their staged nature. Varis provides editor services across stages for client-side code: Starting with an ouput approximation a callgraph with conditional edges is constructed, which

allows navigation in client-site code (HTML, JS and CSS) although it is embedded in server-side code [9, 11].

2.4 Assumptions And Limitations

In addition to the assumptions introduced by symbolic execution, such as symbolic inputs and one-time loop execution, the symbolic execution engine described in previous work [9] does not support advanced language features, including object-oriented programming. Moreover, the range of standard library functions supported are mostly string operations, yet PHP provides a large collection of functions for array operations. All three analyses described in Section 2.3 though have only been applied to systems that are small and not maintained any more. Despite the functional benefits these tools provide to developers, their practicality for large and modern systems with respect to advanced language support has not been evaluated.

3 SCALABLE OUTPUT APPROXIMATION

The main objective of this experience report is to investigate whether we can have a scalable symbolic execution engine as it is the underlying technique for static output approximation in previous work [9–12]. Our vision is a practical symbolic execution, i.e., it computes an accurate approximation for large and modern PHP systems with reasonable time and space consumption.

Also, we are interested in the limitations of scalable symbolic execution for PHP web applications. Limitations may either be conceptional, where symbolic execution of PHP code, in contrast to concrete execution, is inaccurate: For example an include expression evaluated to a symbolic value simply provides to little information to include a script. Or, in spite of infeasible effort required, limitations could allow an accurate symbolic execution. For example, exhaustively exploring all feasible paths in an application may result in an infeasible number of program states to store; nevertheless, this so called state explosion problem may be addressed with additional engineering.

3.1 Towards Scalable Output Approximation?

As our goal is to have a scalable symbolic execution engine, we approach this question by re-implementing the symbolic execution semantics described in previous work [9] and extending it with support for additional language features. We aim to support symbolic execution for a real-world example system of reasonable size that incorporates the missing language features described in Section 2.4 and is still developed and maintained. WordPress is a popular open-source Content Management System (CMS) with around 300.000 lines of code that provides a vivid plug-in environment and is maintained by a large community. For our implementation of a symbolic execution engine for PHP, Oak[2], we have chosen a test-driven approach with continuous integration using regression tests based on WordPress and SchoolMate (see Section 4.1). Oak incorporates the existing execution semantics of [9] and extends it with support for object-oriented orientation (class definition, class instantiation and method invocation) as well as most PHP array operations. A list of the functions provided by the PHP standard library can be found at [?]. Throughout the implementation we did

---

[2]The implementation can be found at github.com/smba/oak.

not achieve high code coverage for WordPress. In the following we present our observations we made thoroughout the implementation regarding issues impeding symbolic execution.

3.2 Experience Report

We identified trade-offs between accuracy and performance that may cause the approximation to be impractical, yet possible. Additionally, we detected limitations impeding symbolic execution conceptually: For certain language features, concrete information rather than symbolic values may be required to actually be executed properly as this defective execution can impact subsequent parts of the program which may become inaccessible and may be missed.

3.2.1 Symbolic vs Concrete Execution. Throughout the implementation, we basically followed the symbolic execution semantics described in previous work [9]. Nevertheless, for additional language features we experienced it to be beneficial to the approximation to execute them as precisely as possible. For instance, WordPress uses a text filter mechanism, where different filter functions are applied sequentially to a string literal. Filters can be linked to a hook using associative arrays so that when the hook is triggered, all filters that are linked to it are applied. As described in Section **??** we execute loops only once due to an possibly indeterminable number of iterations. PHP yet offers more loop constructs: the do-while loop and the foreach loop. While for the do-while loop as well as for the while loop the number of iterations is possibly indeterminable, the foreach loop iterates over all items of a given array whose number of items is known at runtime. For WordPress, we are able to achieve a more accurate approximation and higher code coverage results by concretely executing foreach loops as more server-side code (filter functions in our case) are accessed and executed. In turn, exhaustively executing foreach loops and filter functions concretely is a trade-off between accuracy and performance, as a single hook call site can be called from various contexts since we are exploring multiple execution paths.

An additional aspect on the trade-off between concrete and symbolic execution, or effort and accuracy respectively, is the extent to which program inputs are assumed to be symbolic. As shown in Section **??** introducing symbolic values can be problematic if they are propagated to language features requiring concrete information, such as include expressions or indirect function calls. We manually scaled the number of sources of symbolic inputs by providing static dummy inputs such as a concrete form input, file name or configuration entry. For specific cases, where concrete information is inevitably required, we can achieve higher code coverage and a more accurate approximation, as more code becomes accessible. Nevertheless, while detecting dynamic features can be automated, it requires manual inspection and understanding of the web application to provide appropriate dummy input data. Again, this is a trade-off between the effort to provide input data and accuracy of the approximation and code coverage respectively.

3.2.2 Multi-Target Method Calls. A non-conceptual problem we encountere is related to object-oriented programming. While exploring different execution paths when the control flow is split, a variable can have different values depending on the path condition under which a value has been assigned to it (see the code

example in Figure 1b). Similarly, a variable can point to different objects depending on the path condition as the snippet in Figure 2a illustrates. The variable $person can either point to object Alice or Bob. In addition, for any method invoked on that variable the actual target is ambiguous as there are two, depending on the path condition. In order to have a sound symbolic execution, both method calls need to be executed. Nevertheless, the number of contexts for which a method can be called can grow rapidly, and easily become infeasible to execute exhaustively.

We implemented several heuristic approaches to tame this state space explosion for multi-target method calls, including executing a method only for a subset of target objects, or just a single object. This is a trade-off between accuracy and soundness of the symbolic execution and performance time required.

3.2.3 Dynamic Language Features. Conceptually, we identified two language features for which symbolic execution failed frequently. As stated in Section 2.2, dynamic web applications are staged and client page output is computed at runtime. Symbolic execution has shown to be a useful technique to approximate dynamically generated output in theory, though it is challenged by language features that require concrete information rather than symbolic information at runtime. In particular, this applies to language features that provide access to reusable code of the web application: include expressions and function calls.

First, include expressions in PHP usually require a string value representing the path to the script file or template to include. If this string value is statically provided by a string literal or can be constructed unambiguously, the symbolic execution engine was able to resolve the include expression and include the desired file. However, it is also common to assemble include expressions not only from static string literals, but from several sources including user input information, database query results or configuration files. Figure 2b shows a small example of a dynamic include containing information retrieved from a database. Since any input for a symbolically executed application is symbolic this information can be propagated to include expressions. Also, configuration files represent only one particular configuration or even just default values or placeholders. Hence, symbolic execuion is prone to fail for dynamic includes if symbolic information is contained.

Second, function can either be called directly by statically providing the name of the function (or method), or indirectly. Indirect function calls in PHP are enabled by using built-in functions like `call_user_func`, `call_user_func_array` or simply passing a string value with a direct function call to the `eval` function. The first two functions take as arguments the name of the function two call, and all arguments to that functions either as additional arguments or a single associative array respectively. The `eval` function evaluates a given PHP expression passed as a string literal. Again, we encountered issues similar to dynamic include expressions with indirect function call mechanisms: The example in Figure 2c illustrates two functions `header_serif` and `header_sserif` which both print a HTML headline. Depending on whether the value of the variable $style retrieved from a database is `serif` or `sserif`, a different function is called. Given that the function name is concrete at runtime, the desired function can be called. Otherwise,

```php
1  <?php
2  class Person {
3      function __construct__($name) {
4          $this->name = $name;
5      }
6      function greet($greeting) {
7          echo $greeting . $this->name;
8      }
9  }
10
11 if (...) {
12     $person = new Person("Alice");
13 } else {
14     $person = new Person("Bob");
15 }
16 $person->greet("Good Morning, ");
```

(a) Dynamic dispatch: Multiple targets for method calls

```php
1  <?php
2  define("ROOT", getcwd());
3  define("TEMPLATES", ROOT . "templates/")
4
5  $template = mysql_result(...);
6
7  require_once TEMPLATE . $template;
```

(b) Dynamic Include Resolution

```php
1  <?php
2  function header_serif($title) {
3      $style = "font-family: serif;";
4      echo "<h1 style = '$style'>$title</h1>";
5  }
6  function header_sserif($title) {
7      $style = "font-family: sans-serif;";
8      echo "<h1 style = '$style'>$title</h1>";
9  }
10
11 $font_style = mysql_result(...);
12 call_user_func('header_' . $font_style, 'Title');
```

(c) Function calls by indirect invocation mechanisms

Figure 2: Code Examples of defective code features

if the function name is assembled at runtime and contains symbolic values, no function can be called as the name is ambiguous or unknown. For this example, symbolic execution is not able to determine which function to call indirectly.

3.2.4 Interpreter Customization. We have approached issues we encountered, including multiple-target method calls and foreach loops, in different ways as these represent a trade-off between accuracy of execution and performance. For method calls having ambiguous targets, we implemented different modes of execution, ranging from a complete mode where a method is executed for all possible targets to abstract modes with either a subset of targets or only a single target. In the latter cases where targets are discarded, we assumed symbolic values as method return values. Our symbolic interpreter can be customized to work in the desired mode to scale accuracy or performance. Nevertheless, for WordPress these variations did not have significant impact on the output approximation as symbolic execution was rather impeded by dynamic

features. Also, for foreach loops we are able to toggle between a single iteration mode and a complete iteration mode. Although we were able to increase code coverage by accessing more code, it did not significantly increase our output approximation. Note that these variants have only been tested for WordPress to better understand possible trade-offs or simply achieve any output approximation in a reasonable period of time. We believe based on our experience dynamic features to be the main conceptual limitation impeding code coverage as well as output approximation quality. In the next Section we analyze in detail the impact of dynamic features with concrete foreach loop execution and complete multiple-target execution for method calls.

4 EVALUATION

As already stated in the previous Section 3.2, we encountered several problems during the implementation of our symbolic execution engine for WordPress. The main items of critique of the symbolic execution engine and tools described in previous work [9–12] are the missing support for widely-used language features, and the evaluation only using small-scale systems. Although evaluating a tool for a small systems can be a valuable proof-of-concept, it does not necessarily cover the question of whether a tool is practical on a larger-scale. Therefore in the following we present our methods to evaluate practicality of our symbolic output approximation. Moreover, we investigate to which extent the conceptual limitations encountered with WordPress apply for a wider choice of PHP systems.

4.1 Experiment Setup

We have encountered several issues for WordPress that conceptually limit the scalability of symbolic execution. Since our observations so far only considers one example system an, we evaluate our symbolic execution engine for a wider range of PHP systems with regard to the following questions. The metrics to approach those questions are explained in Sections 4.3 and 4.4.

First, the evaluation of the output approximation using the symbolic execution engine, Symex, in previous work [9] has shown high code coverage for small systems. As our symbolic execution re-implements Symex and extends with support for additional language features, we ask whether we can replicate high coverage results at least on a small-scale. Therefore, our PHP corpus contains four small-scale systems used in previous work [9].

Second, we ask whether system size and usage of modern language features are factors affecting the practicality of our symbolic execution engine. We selected both a range of large-scale and small-scale modern PHP systems to better separate and understand the circumstances under which conceptual limitations might affect practicality. For the selection of large-scale systems we borrow three systems from a case study addressing the feature usage in PHP systems. Some more detailed description of this case study can be found in Section 6. For the small-scale systems we selected five systems from a list of recent Content Management Systems (CMS) written in PHP [1]. For both selections we were bound to select a smaller number of system than we intended to since the parser used in our implementation, Quercus[3], did not support a number of language

---

[3]Quercus footnote

System	Version	Classification	SLOC	#files	#OCs	#includes
AddressBook	8.2.5.2		51,907	239	1009	186
SchoolMate	1.5.4		8,118	65	853	88
TimeClock	1.04		20,800	63	7920	306
WebChess	1.0.0		5,219	28	470	56
Drupal	7.5.0	CMS	52,464	125	3569	749
phpBB	3.1.9		327,371	1,398	3606	206
phpMyAdmin	4.6.3		303,582	871	7103	571
Anchor	0.12.1		15,054	201	987	32
Kirby	XXX	CMS	XXX	XXX	654	23
Automad	XXX	CMS	XXX	XXX	655	8
Monstra	XXX	CMS	XXX	XXX	1934	48
Nibbleblog	XXX	CMS	XXX	XXX	1013	28

Table 1: Corpus of twelve PHP systems. The file count includes files with a .php, .inc, .bit or .module extension.

features. The full list of PHP systems for that we evaluate our symbolic execution engine along with descriptive code metrics can be found in Table 1.

4.2 Measuring Approximation Success

In order to accurately answer the question how good our symbolic execution engine approximates client page output, we would require knowledge of all possible variants to compare our approximation against. Since we do not have any ground truth information alike, we approach this demand by using a heuristic. Rather than having knowledge about all client page output variants we say our approximation is accurate if all string literals (which are embedded in and scattered across server-side code) that may eventually become part of any client page output variant (in the following referred to as *output candidate*) is contained in our symbolic output approximation. It is non-trivial to determine whether a string literal is an output candidate, so we heuristically classify string literals as output candidates if they contain the characters <, >, or both since we expect output to contain HTML tags.

We evaluated this heuristic manually with a sample of 400 string literals randomly selected from our corpus of PHP systems (see Table 1). We measured for our heuristic classifier a precision of 94 percent, and a recall of 50 percent. This means that six percent of the string literals are classified incorrectly as false positives. In turn, the classifier is highly distinctive as 96 percent of the string literals are classified correctly. The recall of 50 percent means that half of the string literals, which we manually classified as output candidates, actually were responsive to the classifier. We attempted to increase recall by looking for further distinctive properties to build a classifier from, but could not do so without decreasing precision. In spite of missing half of the output candidates, we decided to use a simple, yet precise and distinctive classifier.

4.3 Measuring Approximation Accuracy

In Section 4.2 we introduced the definition of output candidates as expected output. For an approximation to be most accurate, it needs to contain all output candidates of a system analyzed. Based on whether output candidates are reached by an execution path and

part of the output approximatiopn define two metrics to approach measuring accuracy.

We measure how much of the expected output was actually processed by the symbolic interpreter. Once a line of code, a statement or expression containing an output candidate is actually an element in an execution path, we define this output candidate as *reached*. We define the metric *reach coverage* as the ratio of output candidates that are reached and the total number of output candidates in the analyzed system. Although a reached output candidate is processed by the symbolic interpreter, it does not guarantee we will see that output candidate in the symbolic output. Therefore, we define the metric *output coverage* as the ratio of output candidates that are contained in the output of the symbolic interpreter and the total number of output candidates in the analyzed system.

The coverage results for our case study are illustrated in Figure 3a. We could replicate high coverage results for the four small-scale systems with a reach coverage and output coverage over 80 percent respectively. For the three modern/large-scale systems we measured poor coverage with reach coverage ranging from 5 to 30 percent and output coverage ranging from 4 to 20 percent. We measured medium to high coverage for the more recent small-scale systems: reach coverage ranging from 43 to 85 percent, output coverage ranging from 40 to 85 percent.

4.4 Understanding Limitations

As we have seen in Section 4.3 both measured coverage metrics were poor for large-scale and small-scale/recent systems. To understand what output candidates we missed and why, we further investigated our approximation results and conducted three more measurements.

4.4.1 What literals did we miss? Our initial approach to understand missed output candidates is to find out whether their surrounding program code was not accessed/accessible, and if so, why. We are interested in those output candidates that are located in HTML files (possibly with nested PHP), or located in a function definition. Plain HTML files do not require any further execution and represent output candidates that just need to be included properly to be part of the output. Aside, output candidates that are
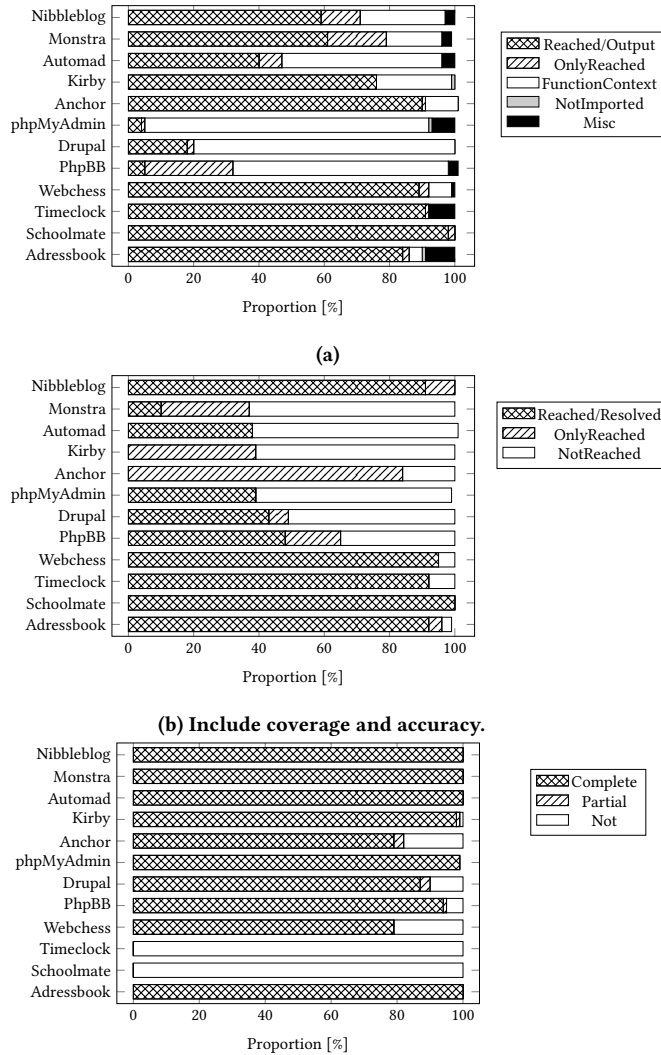
(a)

(b) Include coverage and accuracy.

(c) Distribution dead function output candidates that are par- tially or completely explained due to callback candidates.

Figure 3: Results

part of a function definition are only missed if the corresponding function is never called at runtime.

We started classifying missed output candidates by their string literal context. As the classification statistics in Figure 3a illustrate, for almost all systems tested (except for AddressBook and TimeClock) not-reached output candidates had a function context.

For both cases of context, either inclusion of a script file failed, or a function call failed. If an HTML file is not part of the output, it also is never reached, i.e., included. An include can fail due to various reasons, such as an imprecise evaluation of the include expression returning a symbolic value, or simply a missing file.

For a function to be never reached there are several scenarios: A function is undefined at runtime if the corresponding script file is never included. In turn, if the function is defined at runtime, the

function is dead code if there is no call site for this function, or the function is only called indirectly and the function name could not be resolved accurately. In PHP there are several ways to call a function beside direct call sites.

Given a symbolic value, indirect call mechanisms like callback-commands are likely to fail since target function name and the symbolic value do not match. This also applies for the evaluation of include expressions as concrete values representing include targets can be included, but any include target containing symbolic values is ambiguous. Since it is unlikely that all functions containing missed output candidates are dead code having no call site (direct or indirect) at all, in Section 4.4.3 we further investigate how many functions fail due to inaccurately resolved indirect function calls. Moreover, we evaluate the accuracy of evaluation of include expressions in the next section.

4.4.2 Why did includes fail? From Section 3.2.3 we learn that include expressions can not be evaluated accurately for symbolic values, and from the previous Section 4.4.1 that very little of missed output candidates can be explained by simply missing to include a file. To better understand the impact of include expression resolution throughout analyzing systems, we construct two additional metrics. First, we are interested in how many include expressions throughout the execution we actually reach in a system. Therefore we introduce the metric *reach coverage for include expressions* as the ratio of reached include expressions and the total number of reach expressions in a system to check whether failed includes may have caused subsequent include expressions and so on. Second, we measure how successful include expressions are resolved. Figure 3b illustrates the classification of include expressions for each system: We can either reach an include expression and successfully resolve it, or not. Or, we did not reach an include expression at all. Figure 3b indicates that except for the small-scale systems as well as Anchor and Nibbleblog we did not really reach a great portion of include expressions, although for most systems the resolution success rate was over 80 percent. Note that one failed attempt to include a script file can result in a cascade of missing more include expressions as the include target script file is not included and executed.

4.4.3 Why did function calls fail? Aside from failed or successful includes, this metric describes one scenario of why functions containing not-reached output candidates are never called: The trivial explanation would be that there are no direct calls of those functions, yet PHP offers a number of ways to call a function indirectly, for instance using built-in functions, or simply evaluating strings as PHP code using the `eval` function.

Moreover we need to take into account that one missed function execution can result in missing even more function calls (direct or indirect) and executions. Thus, for functions containing not-reached output candidates we measure whether they can be (1) only, (2) partially or (3) never accessed transitively from functions that have no direct call sites. We refer to those functions (or access points) having no direct call site as callback candidates since the only way to reach them is through indirect mechanisms. Measuring how many not-reached output candidates can be explained by functions that are callback candidates helps to understand whether,

Figure 4: Callgraph analysis: The bottom nodes represent functions containing missed output candidates, the top nodes depict entry points to them.

and if so, indirect call mechanisms and their usage may impact our analysis' code coverage.

The callgraph snippets in Figure 4 illustrate the idea of detecting the functions' access points: The nodes at the top depict callback candidates, the nodes at the bottom with labels X, Y, and Z depict function definitions containing not-reached output candidates. In Figure 4 function X can only be accessed through callback candidates, as for Figure 4 function Y can be partially accessed by callback candidates, whereas function Z is never accessible through callback candidates.

For failed function calls to functions containing not-reached output candidates we measured how many of those functions were only, or partially accessibly through indirect calling mechanisms (see Section 3.2.3). As Figure 3c illustrates, around 80 percent (except for TimeClock) all missed output candidates of each system were only accessible through callback candidates, i.e., functions that are only accessible through indirect calling mechanisms.

5 LESSONS LEARNED

6 RELATED WORK

[6] [8, 13]

7 CONCLUSION

REFERENCES

[1] https://codegeekz.com/12-lightweight-cms-for-building-websites/. (????). Accessed: 2016-09-30.
[2] PHP Usage on GitHub. https://github.com/languages/php. (????).
[3] PHP Usage Statistics. http://w3techs.com/technologies/details/p1-php/all/all. (????). Accessed: 2016-1-12.
[4] TIOBE Index for PHP. . (????).
[5] J. A. Darringer and J. C. King. 1978. Applications of Symbolic Execution to Program Testing. *Computer* 11, 4 (April 1978), 51–60. DOI:http://dx.doi.org/10.1109/C-M.1978.218139
[6] Mark Hills, Paul Klint, and Jurgen Vinju. 2013. An Empirical Study of PHP Feature Usage: A Static Analysis Perspective. In *Proceedings of the 2013 International Symposium on Software Testing and Analysis (ISSTA 2013)*. ACM, New York, NY, USA, 325–335. DOI:http://dx.doi.org/10.1145/2483760.2483786
[7] James C. King. 1976. Symbolic Execution and Program Testing. *Commun. ACM* 19, 7 (July 1976), 385–394. DOI:http://dx.doi.org/10.1145/360248.360252
[8] Yasuhiko Minamide. 2005. Static approximation of dynamically generated web pages. In *Proceedings of the 14th international conference on World Wide Web*. ACM, 432–441.
[9] Hung Viet Nguyen, Christian Kästner, and Tien N. Nguyen. 2014. Building Call Graphs for Embedded Client-side Code in Dynamic Web Applications. In *Proceedings of the 22Nd ACM SIGSOFT International Symposium on Foundations of Software Engineering (FSE 2014)*. ACM, New York, NY, USA, 518–529. DOI:http://dx.doi.org/10.1145/2635868.2635928
[10] Hung Viet Nguyen, Christian Kästner, and Tien N. Nguyen. 2015. Cross-language Program Slicing for Dynamic Web Applications. In *Proceedings of the 2015 10th Joint Meeting on Foundations of Software Engineering (ESEC/FSE 2015)*. ACM, New York, NY, USA, 369–380. DOI:http://dx.doi.org/10.1145/2786805.2786872
[11] Hung Viet Nguyen, Christian Kästner, and Tien N. Nguyen. 2015. Varis: IDE Support for Embedded Client Code in PHP Web Applications. In *Proceedings of the 37th International Conference on Software Engineering - Volume 2 (ICSE '15)*. IEEE Press, Piscataway, NJ, USA, 693–696. http://dl.acm.org/citation.cfm?id=2819009.2819140
[12] Hung Viet Nguyen, Hoan Anh Nguyen, Tung Thanh Nguyen, and Tien N. Nguyen. 2011. Auto-locating and Fix-propagating for HTML Validation Errors to PHP Server-side Code. In *Proceedings of the 2011 26th IEEE/ACM International Conference on Automated Software Engineering (ASE '11)*. IEEE Computer Society, Washington, DC, USA, 13–22. DOI:http://dx.doi.org/10.1109/ASE.2011.6100047
[13] Gary Wassermann and Zhendong Su. 2007. Sound and precise analysis of web applications for injection vulnerabilities. In *ACM Sigplan Notices*, Vol. 42. ACM, 32–41.