

Does It Scale? Static Output Approximation of PHP Web Applications Using Symbolic Execution*

Stefan Mühlbauer

Technische Universität Braunschweig

Christian Kästner

Carnegie Mellon University

Tien N. Nguyen

University of Texas at Dallas

ABSTRACT

Dynamic web applications have become widely popular and are to a large proportion based on the scripting language PHP. Output approximation of web applications enables a range of additional tool support as well as possibilities for vulnerability detection. Unfortunately, recent approximation approaches have only been evaluated for smaller systems.

This paper presents an experience report about the scalability of output approximation using symbolic execution of state-of-the-art PHP web applications. For a symbolic execution engine extended with support for object-oriented programming and arrays, we identified language features and corresponding programming patterns that impede symbolic execution and limit the scalability of this approach. Our findings include: (1) Dynamic features such as functions and includes are prone to fail for certain programming patterns. (2) Expressions containing elements from I/O, databases or files can heavily impede symbolic execution. Our findings provide useful guidelines to design new tools and also to improve the development process of statically analyzable web applications.

CCS CONCEPTS

•**Computer systems organization** Embedded systems; *Redundancy*; Robotics; •**Networks** Network reliability;

KEYWORDS

ACM proceedings, \LaTeX , text tagging

ACM Reference format:

Stefan Mühlbauer, Christian Kästner, and Tien N. Nguyen. 1997. Does It Scale? Static Output Approximation of PHP Web Applications Using Symbolic Execution. In *Proceedings of ACM Woodstock conference, Santa Barbara, California, USA, July 2017 (ISSTA'17)*, 8 pages. DOI: 10.475/123_4

1 INTRODUCTION

With the emerging world wide web, dynamic web applications have become widely popular. Various different implementation techniques have emerged, ranging from technologies for programming languages (JSP, ASP .NET) and web application frameworks for script languages (Ruby On Rails, Django) to languages tailored specifically to the domain of web applications. PHP citephpNET is

a programming language focused on server-side application development. As of 2012, it was used by 78.8 percent of the ten million most popular websites (according to Alexa popularity ranking) [3], ranking 7th on the TIOBE programming community index [4], and was the ranked as the 6th most popular language on GitHub [2].

One common property of all technologies for dynamic web applications is *staged computation* of output: A dynamic web application as a whole consists of both static code, such as scripts, and dynamically generated code, such as responses to HTTP requests. The latter code, although assembled at runtime, may contain client-side parts of the web application, such as JavaScript. So, to study dynamic web applications in its entirety, we need to consider both static as well as dynamic aspects of systems.

Symbolic execution is a static analysis technique between program testing and program proving. For a program, symbolic, i.e., abstract, yet fixed inputs are applied instead of concrete inputs [5, 7]. The underlying concept is to map program input (symbolically) to program output. Different classes of inputs may result in variational output, representing dependencies of output and input. As symbolic execution of a program returns an input-output-mapping for symbolic inputs, this analysis technique helps unfold the *staged* nature of web applications: Since every feasible path is executed all possible output is contained in the symbolic output, and all dynamically generated variational parts of the web application can be analyzed.

Knowledge about different output variants enabled useful analyses for the domain of web applications, such as detecting and locating HTML validation errors [11]. As for developers, from an output model, outlines and callgraphs for easier IDE code navigation [8], or program slices [9] can be computed. Existing work has been integrated in the Varis plugin for the Eclipse IDE [10].

Despite the various use cases for analyses based on an approximated output model, static output approximation for PHP web applications using symbolic execution has so far only been evaluated for smaller systems that are not maintained anymore. The previously mentioned tools though are only practical, if for a given system the symbolic execution engine is scalable, i.e., will also approximate output accurately for larger and more recent systems with good time and space consumption.

To investigate the question, whether we can build practical tools from the the symbolic execution engine, we re-implemented the engine with the specifications of the previous symbolic execution semantics [8] for PHP and additional features, such as object-oriented programming. We evaluated our symbolic execution engine for large-scale and modern PHP systems.

During introduction of new semantics for object-orientation we encountered two main challenges. First, we were bound to introduce more concrete execution features to the symbolic execution

*Produces the permission block, and copyright information

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

ISSTA'17, Santa Barbara, California, USA

© 2016 Copyright held by the owner/author(s). 123-4567-24-567/08/06...\$15.00

DOI: 10.475/123_4

engine, e.g. foreach loops for arrays, as this increased code coverage significantly. Second, we had to make a trade-off between the accuracy of method calls and the scalability of this language feature as a method call for different program state variants can have multiple targets.

Based on our observations and evaluation results, we learned that for dynamic PHP constructs, expressions that are assembled dynamically, in combination with symbolic values symbolic execution can be impractical. For object-orientation and method calls, symbolic execution semantics can be a tightrope walk between state space explosion and imprecision. Moreover, the nature of PHP challenges a practical symbolic execution engine since for it to be accurate, a large number of functions from the standard library needs to be supported.

Our key contributions in this paper include (1) a new tool infrastructure to statically approximate the output for PHP web applications using symbolic execution, and (2) a report of our observations with state-of-the-art web applications as well as an explanation on conceptual limitations for symbolic execution of PHP applications.

2 STATE OF THE ART

Symbolic execution is one of integral analysis techniques used by the tools whose practicality we aim to evaluate. This section recaps the idea of symbolic execution, the tools which make use of it, and both conceptual and language-related limitations.

2.1 Symbolic Execution

Symbolic execution is a static analysis technique that was proposed by James C. King in 1976 [7]. Symbolic execution allows to explore all feasible paths in a program. Based on normal program execution, the execution semantics is extended to handle symbolic values, which are abstract, but fixed; an execution becomes symbolic by introducing symbolic values as program inputs instead of concrete values [5, 7]. These symbolic values are then propagated and used throughout the execution.

Execution starts with a plain path condition, usually a tautology. Whenever the control-flow can be split into different branches, the corresponding path conditions (guards for instance) are conjoined with the previous path condition [7]. Additionally, the conjoined path condition can be checked for satisfiability to exclude infeasible paths. Consequently, for a given program symbolic execution computes a mapping from (symbolic) input values to output values.

2.2 Static Output Approximation for PHP

Symbolic execution is an analysis technique not dedicated to a particular programming language. For dynamic web applications, in our case for web applications written in PHP, it can be used to approximate all possible client page output. Before we present in detail existing tools that aim to ease development and maintenance of PHP web applications, we illustrate how output approximation using symbolic execution is used by starting with an example.

Figure 1 shows two code listings. The snippet of server-side web application code in Figure 1a contains both HTML and PHP. First, the function `make_titles` defined in lines 7 – 13 prints the string `$arg` passed as an argument between zero and an arbitrary number of times. Second, the function is indirectly called using the built-in

```

1  <form>
2    <input type="text" name="name">
3    <input type="submit" value="Submit">
4  </form>
5
6  <?php
7  function make_titles($arg, $n) {
8    $is = shuffle(array(1,2,3));
9    while ($n < $is[0]) {
10     echo "<h1>" . $arg . "</h1><br />";
11     $n = $n + 1;
12   }
13 }
14
15 call_user_func("make_titles", "Headline", 4);
16
17 $greeting = "Hello " . $_POST['name'] . "!";
18 if (!isset($_POST['name'])) {
19   echo "No name entered!";
20 } else {
21   echo $greeting;
22 }
23 ?>
```

(a) Snippet of HTML and PHP server-side web application code.

```

1  <form>
2    <input type="text" name="name">
3    <input type="submit" value="Submit">
4  </form>
5
6  // #repeat n < shuffle(is)
7  <h1>Headline</h1><br />
8  // #endrepeat
9
10 // #if !isset(name)
11 No name entered!
12 // #else
13 Entered α!
14 // #endif
```

(b) Approximation client page output of the server-side web application code in Figure 1a. Variability as well as repetition annotations are represented by preprocessor directives.

Figure 1: Illustration of static output approximation using symbolic execution

function `call_user_func` in line 15 with the string `Headline` and 4 as arguments. Finally, the input from the form in lines 1 – 4 is used: The variable `$greeting` consists of three concatenated strings, `Hello` the value of input field `name`, and an exclamation mark. For this server-side code snippet there are two possible client page output variants: If the input field `name` is empty at runtime, line 19 is executed and `No name entered!` is printed, or, if a name was entered, line 21 is executed and the greeting constructed in line 17 is printed.

The following symbolic execution semantics for PHP features based on [8] illustrate the necessary modifications made to handle the extent of ambiguity that comes with handling symbolic values. In addition to normal output, the output approximation contains preprocessor directives to express repeatable output as well as output variants. All approximated client page output corresponding to Figure 1a consisting of two different variants is illustrated in Figure 1b. First, loops in the previous execution semantics are only executed

once instead of an arbitrary number of times. Since a loop condition can contain symbolic values, it might not be feasible to determine when a loop terminates. Hence, any output constructed during the single loop iteration is highlighted (`repeat ... endrepeat`) to possibly be repeated an arbitrary number of times. The output in line 7 is marked to be repeated an arbitrary number of times, as the number of loop iterations cannot be determined statically. Second, a similar procedure applies to recursive functions. Once called, a recursive function is also only executed to recursion depth one per call, and any subsequent invocation of that function returns a symbolic value. This, again, is due to the difficulty to determine the exact recursion depth similar to the number of loop iterations. Third, since the symbolic execution engine applies a static analysis, any form of user interaction providing input data as well as any form of input referred to from the deployment context of the web application, i.e. configuration files or databases, is substituted by a symbolic value. The context variable `$_POST` represents input submitted to the server and, hence, is symbolic as every program input is per definition substituted by symbolic values. Finally, the output in lines 11 and 13 respectively are marked as alternatives (`if ... else ... endif`), as symbolic execution of lines 18 – 22 in Figure 1a will explore both different branches.

2.3 Existing Tool Support for PHP

All tools proposed so far leverage a representation of all possible HTML client page output. Since any variant can be analyzed in particular, this enables tool support addressing the web application as a whole in spite of its staged nature, where parts of the web application (client-side output) are generated dynamically. Based on the approximated output representation subsequent analyses can be conducted, which have been implemented in the following tools:

- *PhpSync*: Using all variants of HTML client page output, every single one can be statically checked for Markup Validity, i.e., if the web page conforms to syntactical specifications for HTML and other client-side languages. After validation errors are detected, auto-fixes are provided. Otherwise, the validation error can be traced back to source code responsible for the defect, which then can be refactored manually [11].
- *WebSlice*: Program slices enable to extract and understand the impact of changes in an application. To consider all of a dynamic web application for program slices, client page output needs to be taken into account. WebSlice combines PHP data-flow information with an output approximation to enable program slices across different languages [9].
- *Varis*: Editor services such as “jump to declaration” are nontrivial for dynamic web applications due to their staged nature. Varis provides editor services across stages for client-side code: Starting with an output approximation a callgraph with conditional edges is constructed, which allows navigation in client-site code (HTML, JS and CSS) although it is embedded in server-side code [8, 10].

2.4 Assumptions And Limitations

In addition to the limitations introduced by symbolic execution, such as symbolic inputs and one-time loop execution, the symbolic execution described in [8] does not support advanced language features, including object-oriented programming. Moreover, the range of standard library functions supported are mostly string operations, yet PHP provides a large collection of functions for array operations. All three tools described in Section 2.3 though have only been applied to systems that are small and not maintained any more. Despite the functional benefits these tools provide to developers, their practicality for large and modern systems with respect to advanced language support has not been evaluated.

3 SCALABLE OUTPUT APPROXIMATION

The main objective of this experience report is to investigate whether we can have built practical tools that are on output approximations like the tools in Section 2.3. In particular, we want to find out whether symbolic execution as the underlying technique for static output approximation is scalable, i.e., it computes an accurate approximation for large and modern PHP systems with reasonable time and space consumption.

Also, we are interested in the limitations of scalable symbolic execution for PHP web applications. Limitations may either be conceptual, where symbolic execution of PHP code, in contrast to concrete execution, is inaccurate: For example an include expression evaluated to a symbolic value simply provides to little information to include a script. Or, in spite of infeasible effort required, limitations could allow an accurate symbolic execution. For example, exhaustively exploring all feasible paths in an application may result in an infeasible number of program states to store; nevertheless, this so called state explosion problem may be addressed with additional engineering.

3.1 Towards Scalable Output Approximation?

As our goal is to have practical tools support, we approach this question by re-implementing the symbolic execution semantics described in [8] and extending it with support for additional language features. We choose to aim to support symbolic execution for a real-world example system of reasonable size that incorporates the missing language features described in Section 2.4 and is still developed and maintained: WordPress is a popular open-source Content Management System (CMS) with around 300k lines that provides a vivid plug-in environment and is maintained by a large community. For our implementation of a symbolic execution engine for PHP, Oak, we have chosen a test-driven approach with continuous integration using regression tests based on WordPress and SchoolMate (see Section 4.2). Oak incorporates the existing execution semantics of [8] and extends it with support for object-oriented orientation (class definition, class instantiation and method invocation) as well as PHP array operations. A list of the functions provided by the PHP standard library can be found at [?]. Throughout the implementation we did not achieve high code coverage for WordPress, which was the main motivation to further investigate limitations in the wild.

3.2 Things We Learned The Hard Way

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

3.2.1 Symbolic vs Concrete Execution. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

3.2.2 Multi-Target Method Calls. The first non-conceptual problem we encountered, in particular for WordPress is related to object-oriented programming. While exploring different execution paths when the control flow is split, a variable can have different values depending on the path condition under which a value has been assigned to it (see the code example in Figure 1b). Similar, a variable can point to different objects depending on the path condition as the snippet in Figure 2a illustrates. The variable `$person` can either point to object `Alice` or `Bob`. In addition, for any method invoked on that variable, the actual target is ambiguous as there are two, depending on the path condition. In order to have a sound symbolic execution, both method calls need to be executed. Nevertheless, the number of contexts for a method can grow rapidly, and easily become infeasible to execute.

We implemented several heuristic approaches to tame this state space explosion for multi-target method calls, including executing a method only for a subset or just a single object. This trade-off between accuracy and soundness of the symbolic execution and performance time required.

3.2.3 Dynamic Language Features. Conceptually, symbolic execution failed for two types of language features. As stated in Section 2.2, dynamic web applications are staged and client page output is computed at runtime. Symbolic execution hand has shown to be a useful technique to approximate dynamically generated output, it is challenged by language features that require concrete information rather than symbolic information at runtime. In particular, this applies to language features that provide access to reusable code of the web application: include expressions and function calls.

First, include expressions in PHP usually require a string value representing the path to the script file or template to include. If this string value is statically provided by a string literal or can be constructed unambiguously, the symbolic execution engine was able to resolve the include expression and include the desired file. However, it is also common to assemble include expressions not only static string literals, but from several sources including user input information, database query results or configuration files. Figure 2b shows a small example of a dynamic include containing information retrieved from a database. Since any input for a symbolically executed application is symbolic this information leads to include expressions containing symbolic values. Also, configuration files represent only one particular configuration or even just default values or placeholders. Hence, symbolic execution is prone to fail for dynamic includes if symbolic information is contained. Second, function calls can either be direct by statically providing the name of the function (or method), or indirect. Indirect function calls in PHP are enabled by using built-in functions like `call_user_func`, `call_user_func_array` or simply passing a string value with a direct function call to the `eval` function. The first two functions take as arguments the name of the function to call, and all arguments to that functions either as additional arguments or array respectively. The `eval` function evaluates a given PHP expression passed as a string literal. Again, we encountered similar issues to dynamic include expressions with indirect function call mechanisms. The example in Figure 2c illustrates two functions `header_serif` and `header_sserif` which both print a HTML headline. Depending on whether the value of the variable `$style` retrieved from a database is `serif` or `sserif`, a different function is called. Given that the function name is concrete at runtime, the desired function can be called. Otherwise, if the function name is assembled at runtime and contains symbolic values, no function can be called as the name is ambiguous or unknown. For this example, symbolic execution is not able to determine which function to call indirectly.

3.3 Things We Tried...

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero

```

1 <?php
2 class Person {
3     function __construct($name) {
4         $this->name = $name;
5     }
6     function greet($greeting) {
7         echo $greeting . $this->name;
8     }
9 }
10
11 if (...) {
12     $person = new Person("Alice");
13 } else {
14     $person = new Person("Bob");
15 }
16 $person->greet("Good Morning, ");

```

(a) Dynamic dispatch: Multiple targets for method calls

```

1 <?php
2 define("ROOT", getcwd());
3 define("TEMPLATES", ROOT . "templates/")
4
5 $template = mysql_result(...);
6
7 require_once TEMPLATE . $template;

```

(b) Dynamic Include Resolution

```

1 <?php
2 function header_serif($title) {
3     $style = "font-family: serif;";
4     echo "<h1 style = '$style'>$title</h1>";
5 }
6 function header_sserif($title) {
7     $style = "font-family: sans-serif;";
8     echo "<h1 style = '$style'>$title</h1>";
9 }

```

(c) Function calls by indirect invocation mechanisms

Figure 2: Code Examples of defective code features

ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

4 EVALUATION

As already stated in the previous section, we encountered several problems during the development of the symbolic interpreter. Therefore this section describes our methods to evaluate practicality of the output approximation, and possible explanations.

4.1 Measuring Approximation Success

For the evaluation of our output approximation we require a ground truth output model to compare our approximation against. Since output literals are embedded in server-side code and scattered, it is nontrivial to determine whether a string literal is part of the output or not. So, in the absence of ground truth for the approximated output we are bound to choose a *heuristic approach* in order to measure the accuracy of our approximation. We define a string literal to be an *output candidate* if there exists an execution path

reaching the string literal and eventually passing it to an output-generating statement (e.g., `echo` or `print`).

We approximate output candidates as those string literals containing the characters `<`, `>`, or both since we expect output to contain HTML tags. We evaluated this heuristic manually with a sample of 400 string literals randomly selected from the entire corpus. We measured for our heuristic classifier a precision of 94 percent, and a recall of 50 percent. This means that six percent of the string literals are classified incorrectly as false positives. In turn, the classifier is highly distinctive as 96 percent of the string literals are classified correctly. The recall of 50 percent means that half of the string literals, which we manually classified as output, actually were responsive to the classifier. We attempted to increase recall by looking for further distinctive properties to build a classifier from, but could not do so without decreasing precision. In spite of missing half of the output candidates, we decided to use a simple, yet distinctive classifier.

4.2 Experiment Setup

For the evaluation we symbolically executed a selection of PHP systems. We selected a corpus of twelve PHP systems with regard to system size and recency as the corpus of the case study for [8] only contained small-scale systems that are not maintained any more. The full list of PHP systems is shown in Table 1. Our selection of PHP systems includes

- four small-scale systems that selected from the previously mentioned case study in order to compare our results with the previous symbolic execution engine [8],
- three recent large-scale systems selected from the case study corpus of [6], and
- four small-scale systems selected from a list of recent Content Management Systems [1].

As for for the last two cases, we limited our selection since the parser we used did not support all language features used. For the experimental evaluation we symbolically executed per system each script file or entry point (files with a `.php`, `.inc`, `.bit` or `.module` extension). All following measurements are cumulated per system for each entry point.

4.3 Measuring Approximation Accuracy

In Section 4.1 we introduced the definition of output candidates as expected output. As a most accurate approximation contains all output candidates, we define two metrics to measure accuracy of our approximation.

First, we measure how much of the expected output was actually processed by the symbolic interpreter. Once a line of code, a statement or expression containing an output candidate is actually an element in an execution path, we define this output candidate as reached. We define the metric *reach coverage* as the ratio of output candidates that are reached and the total number of output candidates in the analyzed system. Although a reached output candidate is processed by the symbolic interpreter, it does not guarantee we will see that output candidate in the symbolic output. Second, we define the metric *output coverage* as the ratio of output candidates that are contained in the output of the symbolic interpreter and the total number of output candidates in the analyzed system. For

System	Version	Classification	SLOC	#files	#OCs	#includes
AddressBook	8.2.5.2		51,907	239	1009	186
SchoolMate	1.5.4		8,118	65	853	88
TimeClock	1.04		20,800	63	7920	306
WebChess	1.0.0		5,219	28	470	56
Drupal	7.5.0	CMS	52,464	125	3569	749
phpBB	3.1.9		327,371	1,398	3606	206
phpMyAdmin	4.6.3		303,582	871	7103	571
Anchor	0.12.1		15,054	201	987	32
Kirby	XXX	CMS	XXX	XXX	654	23
Automad	XXX	CMS	XXX	XXX	655	8
Monstra	XXX	CMS	XXX	XXX	1934	48
Nibbleblog	XXX	CMS	XXX	XXX	1013	28

Table 1: Corpus of twelve PHP systems. The file count includes files with a .php, .inc, .bit or .module extension.

an imprecise approximation, we may see loss of output candidate information resulting in a output coverage lower than the reach coverage.

The coverage results are illustrated in Figure 4a. We could replicate high coverage results for the four small-scale systems with a reach coverage and output coverage over 80 percent respectively. For the three modern/large-scale systems we measured poor coverage with reach coverage ranging from 5 to 30 percent and output coverage ranging from 4 to 20 percent. We measured medium to high coverage for the more recent small-scale systems: reach coverage ranging from 43 to 85 percent, output coverage ranging from 40 to 85 percent.

4.4 Understanding Limitations

As we have seen in Section 4.3 both measured coverage metrics were poor for large-scale and small-scale/recent systems. To understand what output candidates we missed and why, we further investigated our approximation results and conducted three more measurements.

4.4.1 What literals did we miss? Our initial approach to understand missed output candidates is to find out whether their surrounding program code was not accessed/accessible, and if so, why. Of special interest are those output candidates that are located in HTML files (possibly with nested PHP), or located in a function definition. Plain HTML files do not require any further execution (of course, unless PHP scripts are embedded) and represent output candidates that just need to be included properly to be part of the output. Aside, output candidates that are part of a function definition are only missed if the corresponding function is never called at runtime.

We started classifying missed output candidates by their string literal context. As the classification statistics in Figure 4a illustrate, for almost all systems (except for AddressBook and TimeClock) not-reached output candidates had a function context. Note that SchoolMate was excluded from the diagram since our analysis reached all output candidates for that system.

4.4.2 Inaccessible Dynamic Features. For both cases of context, either inclusion of a script file failed, or a function call failed. If an HTML file is not part of the output, it also is never reached, i.e., included. An include can fail due to various reasons, such as an imprecise evaluation of the include expression returning a symbolic value, or simply a missing file.

For a function to be never reached there are several scenarios: A function is undefined at runtime if the corresponding script file is never included. In turn, if the function is defined at runtime the function can either be dead code if there is no call site for this function, or the function call failed. In PHP there are several ways to call a function beside direct call sites. The language offers indirect call mechanisms like callback-commands or an evaluation function that parses and evaluates PHP source code as strings.

Given a symbolic value, indirect call mechanisms like callback-commands are likely to fail since target function name and the symbolic value do not match. This also applies for the evaluation of include expressions as concrete values representing include targets can be included, but any include target containing symbolic values is ambiguous.

4.4.3 Why did includes fail? To better understand include expression evaluation, we use another metric to approach accuracy from a different angle: We measure (1) the ratio of reached include expressions and the total number of include expressions in a system respectively (coverage of include expressions), and (2) the ratio of successfully resolved include expressions and reached include expressions in a system (resolution success rate). For a most accurate result all include expressions are reached and resolved.

For include expressions we sampled failed attempts and manually evaluated the cause of failure. We identified for (1) that all failed attempts were dynamic include expressions, i.e., those which were resolved at runtime and were assembled from fragments of environment-dependent information including database retrievals, user input, or files. Consequently, the resolved include expressions contained symbolic values which did not represent identifiable script files.

Finally, we measured for include expressions the overall reach coverage and the overall resolution success rate per system. As




Figure 3: Callgraph analysis: The bottom nodes represent functions containing missed output candidates, the top nodes depict entry points to them.

output candidates we measured how many of those functions were only, or partially accessible through indirect calling mechanisms (see Section 4.4.2). As Figure 4c illustrates, around 80 percent (except for TimeClock) all missed output candidates of each system were only accessible through callback candidates, i.e., functions that are only accessible through indirect calling mechanisms.

illustrated in Figure 4b, except for the small-scale systems and two recent small-scale systems we did not even really reach a great portion of include expressions, although for most systems the resolution success rate was over 80 percent. Nevertheless, one failed attempt to include a script file can result in a cascade of missing more include expressions as the include target script is not executed.

4.4.4 Why did function calls fail? Aside from failed or successful includes, this metric describes one scenario of why functions containing not-reached output candidates are never called: The trivial explanation would be that there are no direct calls of those functions, yet PHP offers a number of ways to call a function indirectly, for instance using built-in functions, or simply evaluating strings as PHP code using the `eval` function.

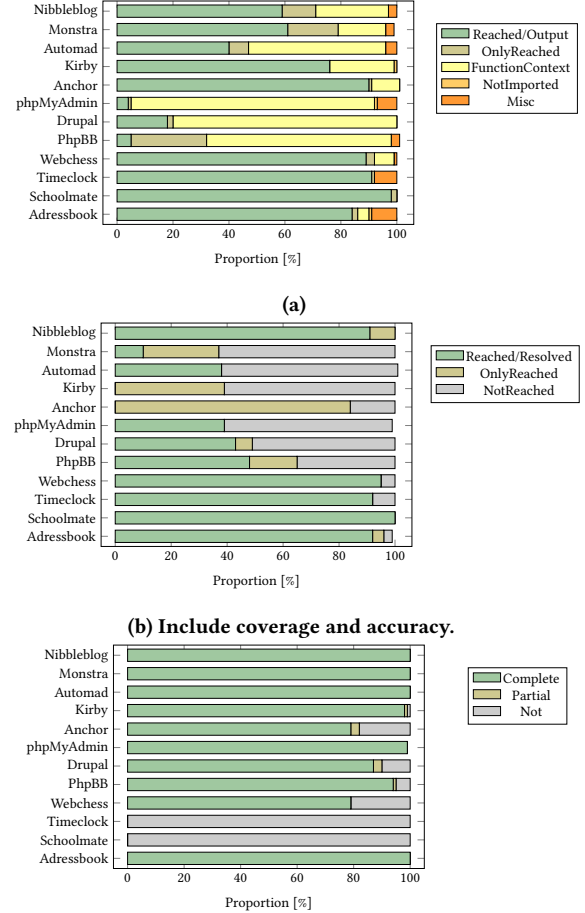
Moreover we need to take into account that one missed function execution can result in missing even more function calls (direct or indirect) and executions. Thus, for functions containing not-reached output candidates we measure whether they can be (1) only, (2) partially or (3) never accessed transitively from functions that have no direct call sites. We refer to those functions (or access points) having no direct call site as callback candidates since the only way to reach them is through indirect mechanisms. Measuring how many not-reached output candidates can be explained by functions that are callback candidates helps to understand whether, and if so, indirect call mechanisms and their usage may impact our analysis’ code coverage.

The callgraph snippets in Figure 3 illustrate the idea of detecting the functions’ access points: The nodes at the top depict callback candidates, the nodes at the bottom with labels X, Y, and Z depict function definitions containing not-reached output candidates. In Figure 3 function X can only be accessed through callback candidates, as for Figure 3 function Y can be partially accessed by callback candidates, whereas function Z is never accessible through callback candidates.

Like for include expressions, we again sampled dead function cases and identified the lack of direct call sites to those functions. As described in Section 4.4.2, functions despite having no direct call sites were called through indirect call features. We identified dynamically assembled expressions (function names) to be responsible for failed attempts of calling a function indirectly as the resolved

function name eventually contained symbolic values. These expressions, as well as include expressions, contained information dependent on the system environment.

For failed function calls to functions containing not-reached



(c) Distribution dead function output candidates that are partially or completely explained due to callback candidates.

Figure 4: Results

5 LESSONS LEARNED

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

6 RELATED WORK

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit

amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

REFERENCES

- [1] <https://codegeekz.com/12-lightweight-cms-for-building-websites/>. (????). Accessed: 2016-09-30.
- [2] PHP Usage on GitHub. <https://github.com/languages/php>. (????).
- [3] PHP Usage Statistics. <http://w3techs.com/technologies/details/p1-php/all/all>. (????). Accessed: 2016-1-12.
- [4] TIOBE Index for PHP. . (????).
- [5] J. A. Darringer and J. C. King. 1978. Applications of Symbolic Execution to Program Testing. *Computer* 11, 4 (April 1978), 51–60. DOI:<http://dx.doi.org/10.1109/C-M.1978.218139>
- [6] Mark Hills, Paul Klint, and Jurgen Vinju. 2013. An Empirical Study of PHP Feature Usage: A Static Analysis Perspective. In *Proceedings of the 2013 International Symposium on Software Testing and Analysis (ISSTA 2013)*. ACM, New York, NY, USA, 325–335. DOI:<http://dx.doi.org/10.1145/2483760.2483786>
- [7] James C. King. 1976. Symbolic Execution and Program Testing. *Commun. ACM* 19, 7 (July 1976), 385–394. DOI:<http://dx.doi.org/10.1145/360248.360252>
- [8] Hung Viet Nguyen, Christian Kästner, and Tien N. Nguyen. 2014. Building Call Graphs for Embedded Client-side Code in Dynamic Web Applications. In *Proceedings of the 22Nd ACM SIGSOFT International Symposium on Foundations of Software Engineering (FSE 2014)*. ACM, New York, NY, USA, 518–529. DOI:<http://dx.doi.org/10.1145/2635868.2635928>
- [9] Hung Viet Nguyen, Christian Kästner, and Tien N. Nguyen. 2015. Cross-language Program Slicing for Dynamic Web Applications. In *Proceedings of the 2015 10th Joint Meeting on Foundations of Software Engineering (ESEC/FSE 2015)*. ACM, New York, NY, USA, 369–380. DOI:<http://dx.doi.org/10.1145/2786805.2786872>
- [10] Hung Viet Nguyen, Christian Kästner, and Tien N. Nguyen. 2015. Varis: IDE Support for Embedded Client Code in PHP Web Applications. In *Proceedings of the 37th International Conference on Software Engineering - Volume 2 (ICSE '15)*. IEEE Press, Piscataway, NJ, USA, 693–696. <http://dl.acm.org/citation.cfm?id=2819009.2819140>
- [11] Hung Viet Nguyen, Hoan Anh Nguyen, Tung Thanh Nguyen, and Tien N. Nguyen. 2011. Auto-locating and Fix-propagating for HTML Validation Errors to PHP Server-side Code. In *Proceedings of the 2011 26th IEEE/ACM International Conference on Automated Software Engineering (ASE '11)*. IEEE Computer Society, Washington, DC, USA, 13–22. DOI:<http://dx.doi.org/10.1109/ASE.2011.6100047>